

The pipeline as it stands, 21 September 2026

Limb-segment volume from phone photos: about eight photos of a limb wearing two green bands, standing next to a 10 cm ArUco cube, give the volume of the limb between the bands. This document describes what the code on `keng-branch` does today (version 4 with the limb skeleton step), stage by stage. It replaces the parts of `pipeline.md` that have gone out of date: the photos are now padded, not cropped; the cut is made on the mesh in Stage 5; the cube is 10 cm, not 14 cm.

At a glance

```
photos (about 8, phone, 9:16)
[0] Framing      find cube, limb and bands; refuse unusable photos
[1] VGGT        whole photo, padded to 518 px; one 3-D point per pixel
[2] Point cloud  keep confident pixels, drop the white padding
[3] Clean       separate cube and limb, stacked MLS, level to the floor,
               find the band planes, limb skeleton, close the ends
[4] Mesh        Poisson surface; alpha shape if Poisson is not one solid
[5] Watertight+cut repair to a closed solid, then cut between the bands
[6] Volume      scale from the cube's volume; limb volume, girths
```

stage	time on test6
1 VGGT	17 s
2 point cloud	under 1 s
3 clean	8 s
4 mesh	16 s
5 watertight and cut	4 s
6 volume	under 1 s
whole run, including Stage 0 and loading the models	85 s

Current accuracy, volume between the bands against water or tape:

capture	truth (cm ³)	pipeline (cm ³)	error
test6	1398.6 tape	1485.2	+6.2%
test5	2070 water	2167.6	+4.7%
0_right	1830 water	1743.5	-4.7%
1_left	1600 water	1591.1	-0.6%
6_left	2800 water	2647.2	-5.5%
2_left	1050 water	1180.2	+12.4%
mean error, ignoring sign			5.7%
mean error, with sign			+2.1%

These are Stages 3–6 rerun with the skeleton step on, from the padded runs' Stage 0–2 output. The previous version, which cropped each photo, averaged 29%.

Stage 0 — Framing

Code: `pipeline/stages/prep.py, prepare_frames()`

For every photo:

- **Cube:** found by its ArUco markers (`DICT_5X5_250`), each marker's quad grown to the whole cube face, joined with a GroundingDINO box.
- **Limb and bands:** GroundingDINO boxes and SAM masks. Every band box that sits on the limb is kept, up to two.
- **Band colour:** measured from the photos and handed to Stage 3, which uses it only when the band and the skin are different enough in colour.
- **Gate:** a photo with no cube at all, or with nothing detected, or that cannot be read, is rejected and stops the run (unless `--continue-on-rejected`). A missing or clipped band is only a warning.

No crop. Stage 0 used to cut each photo to a square that slid up or down to fit the cube and bands. That moved the photo's optical centre 19–74 px away from the middle of the square, while VGGT assumes the centre is the middle, and the whole reconstructed scene came out distorted. Stage 0 now detects and gates only, and hands the whole photo on (`FRAME_FIT = "pad"`). This one change took the mean error from 29% to about 6%.

Output: `framing.json` (band count, band colour), `manifest.json` (per photo: boxes, verdict).

Stage 1 — VGGT

Code: `pipeline/stages/inference.py, run_inference()`

One pass of VGGT-1B over all the photos together. Each photo is shrunk to fit a 518 px square, whole, with white padding above and below (VGGT's own `pad` mode). For every pixel of every photo VGGT predicts a 3-D point in one shared scene, and how confident it is. Those points (`world_points`) are the only geometry the rest of the pipeline uses.

- **Checkpoint:** `facebook/VGGT-1B-Commercial`, the commercial licence (`VGGT_USE_COMMERCIAL = True`). It is gated on Hugging Face; without a login the loader falls back, loudly, to the non-commercial weights.
- **Resolution:** fixed at 518. Larger inputs (1022) were tested and are much worse; the model was trained at 518.

Output: `predictions.npz`

Stage 2 — Point cloud

Code: `pipeline/stages/pointcloud.py, export_ply()`

1. Keep pixels above the 45th confidence percentile (`--conf_thres 45`), with an absolute floor for photos where most pixels sit at the minimum.
2. Drop every pixel of the white padding: VGGT still predicts a point there, and it has no geometry.
3. Remove spatial outliers.

TSDF fusion of the depth maps is available as an alternative (`POINTCLOUD_METHOD=tsdf`) but is off: it was not more accurate.

Output: `points.ply`

Stage 3 — Clean

Code: `pipeline/stages/clean.py, clean_and_extract()`

Separate the objects (in VGGT's own coordinates)

1. Outlier removal, then thinning on a small voxel grid, for speed.

2. Remove the floor plane, so the objects do not join through it.
3. DBSCAN into clusters; the cube is the one that looks most like a cube with black and white faces, the limb is the other.
4. Find the band planes two ways: by colour on the dense limb, and by projecting Stage 0's band boxes through VGGT's points into 3-D. The projected planes only add; where both find a band, the colour plane is used and the projected one is kept as its backup.

Clean each object on its own


5. Ghost-voxel downsample (voxel 2.8 times the point spacing) and a filter that drops points whose normal disagrees with their neighbours.
6. **Stacked MLS.** VGGT gives the skin twice, as two sheets a few mm apart, because different photos place it slightly differently. MLS moves each point onto a smooth patch fitted to its neighbours inside a ball: first a ball of 4 point spacings (1.1 cm on `test6`), then, for the limb only, 16 spacings (2.6 cm), wide enough to reach both sheets and merge them into one. See `docs/pipeline/mls_explained.md`.

Level and finish

7. Fit the floor, rotate so it is flat and height is straight up, and drop anything below it.
8. Gate the band planes: a colour plane must sit at least one cube height above the floor, and every plane must lie within 35° of square to the limb. With two planes left, the measurement is the segment between them.
9. **Limb skeleton** (on by default). Between the bands, the limb is cut into 60 horizontal chunks (0.53 cm tall on `test6`), each chunk's points are seen from above as a ring, and one smooth curve is fitted around the chunk's centre. Points far outside it (spurs) are moved in, points far inside (dents) moved out, and empty arcs are filled on the curve. Section below; full account in `docs/pipeline/limb_skeleton.md`.
10. Close the limb: sweep its bottom down to the floor, cap the bottom, cap the open top. The limb is **not cut here**.


Output: `objects/leg.ply`, `objects/box.ply`, and the band planes in `debug/cutting_line_levelled.json`

The limb skeleton step



- **Chunks:** 0% is the lower band's height, 100% the upper band's; 60 chunks run from -8% to 108%. Inside a chunk, height is ignored.
- **Skeleton:** each chunk's centre from a circle fit, smoothed up the leg.

- **Curve:** walking once around the centre, the distance to the skin rises and falls; the curve is the average distance plus up to six waves (two per turn makes an oval, three a rounded triangle). Higher waves are held back, and the fit is repeated three times without the points furthest from it, so a spur cannot pull the curve toward itself.
- **Unusual points:** further than 2% of the radius (about 1 mm) from the curve. They move only to the edge of that band, so real skin texture stays.



Stage 4 — Mesh

Code: `pipeline/stages/reconstruct.py`

Poisson surface reconstruction for both the cube and the limb. After each, the same repair Stage 5 will run is tried and the result's Euler characteristic checked: a single closed solid has $\chi = 2$. If Poisson's mesh is not (a closed surface with tunnels still counts as "watertight", so that test alone is not enough), that object alone is rebuilt with an alpha shape, whose ladder of sizes is chosen to give $\chi = 2$.

Output: `mesh/box_recon.ply`, `mesh/leg_recon.ply`

Stage 5 — Watertight and cut

Code: `pipeline/stages/watertight.py`, `pipeline/core/meshcut.py`

1. Repair each mesh to a closed solid with PyMeshFix (skipped when it is already closed). The uncut limb is published as `leg_no_cut.ply`.
2. **Cut the limb** at the band planes: two planes keep what lies between them, one keeps what lies below it. Slicing a closed solid gives the flat end faces directly, with no reconstruction step after the cut that could round them off. This is also the surface a person can move the cut on in the web review.

The cut can instead be made on the point cloud in Stage 3 (`CUT_STAGE=cloud`). On the six captures the two agree within 8 cm³ on five, and the cloud cut is further from water on the sixth, so the mesh cut stays.

Output: `mesh/leg_no_cut.ply`, `mesh/leg_cut.ply`, `mesh/box.ply`

Stage 6 — Volume

Code: `pipeline/stages/volume.py`, `compute_volumes()`


- **Scale from the cube's volume.** The cube is $10 \times 10 \times 10$ cm (`REFERENCE_REAL_SIZE_CM`), so its mesh must be 1000 cm³; the ratio gives cm³ per scene unit, and its cube root cm per unit. The printed 5 cm markers were tried as a second scale and dropped: VGGT does not place their corners well enough, and a ruler agreed with the cube.
- **Cube check:** the cube should fill 0.87–0.89 of its own oriented box on a sound capture; below 0.83

the run warns that the reconstruction is suspect.

- **Limb volume:** the exact enclosed volume of the closed `leg_cut.ply`. Voxel counting is only a fallback for a mesh that is not closed, and warns when used.
- **Girth at each band:** the limb's points in a ± 4 mm slab at each cut plane, an ellipse fitted to them, and its perimeter. This is the one number a tape measure can check without water.

Output: `volumes.csv`; the run's final meshes in the output folder.

How it got here



version	date	what changed	mean error
v1	main	VGGT's centre crop, one band, cut on the cloud	not comparable
v2	12–24 Aug	Stage 0 framing; Poisson with alpha fallback	—
v3	31 Aug	cut moved to the mesh in Stage 5; two bands	31% on five captures; 2_left lost a band (+222%)
v4	18–21 Sep	second MLS pass, ghost-voxel spacing fix, projected band planes as backup (29% while still cropping), then pad instead of crop , marker-scale check removed, limb skeleton	5.7%

What is still open

- **Length on `test6`.** The leg reads 28.87 cm between the band centres against a 27.5 cm ruler (+5%), most of `test6`'s error. Whether the ruler was read between band centres or inner edges is not yet known.
- **2_left at +12%.** Its band girths match the tape almost exactly, so the extra is in length or in the water truth.
- **How repeatable the truth and the pipeline are.** Three captures and three water measurements of one limb would say how much of the 5–6% is noise.
- **Six captures** are too few to separate a +2% bias from scatter.

Running it

```
python run.py -i inputs/test6 # the whole pipeline
python stagerun.py 3-6 --name test6_try # rerun Stages 3-6 on work/test6_try
LIMB_SKELETON=off python run.py -i inputs/test6 # without the skeleton step
FRAME_FIT=crop python run.py -i inputs/test6 # the old cropped input
CUT_STAGE=cloud python run.py -i inputs/test6 # cut the cloud in Stage 3
```

Outputs land in `output/` (or `--output_dir`), with every stage's files under `for_debug/00_prep` to `for_debug/06_volume`; `stagerun.py` writes to `work/<name>/` instead.